

Breaking the Ice with OpenStack Neutron

Muneeb Ahmad

About me..

Senior Design Engineer at PLUMgrid Inc with over three years of experience in implementing software-defined networking (SDN), high availability, distributed systems and APIs. An active member of the OpenStack community, with contributions in the networking ecosystem, focused primarily on Neutron project in OpenStack.



muneeb@plumgrid.com

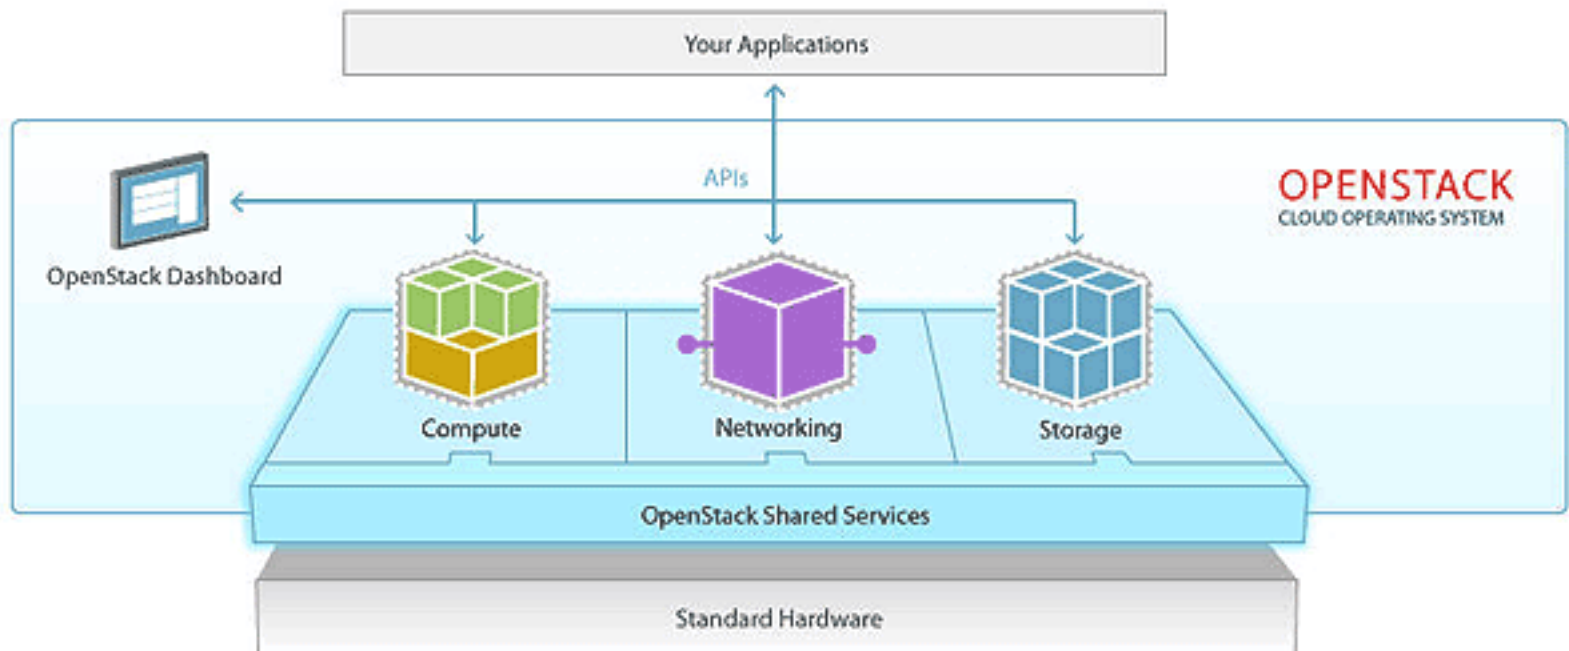


[@smuneebahmed](https://twitter.com/smuneebahmed)



What is OpenStack?


- **OpenStack** is a cloud management system that controls large pools of compute, storage, and networking resources throughout a datacenter, all managed through a dashboard that gives administrators control while empowering their users to provision resources through a web interface.



OpenStack Core Services



Compute (Nova) provides virtual servers on demand. Compute resources are accessible via APIs for developers building cloud applications and via web interfaces for admins and users. The compute architecture is designed to scale horizontally on standard hardware.

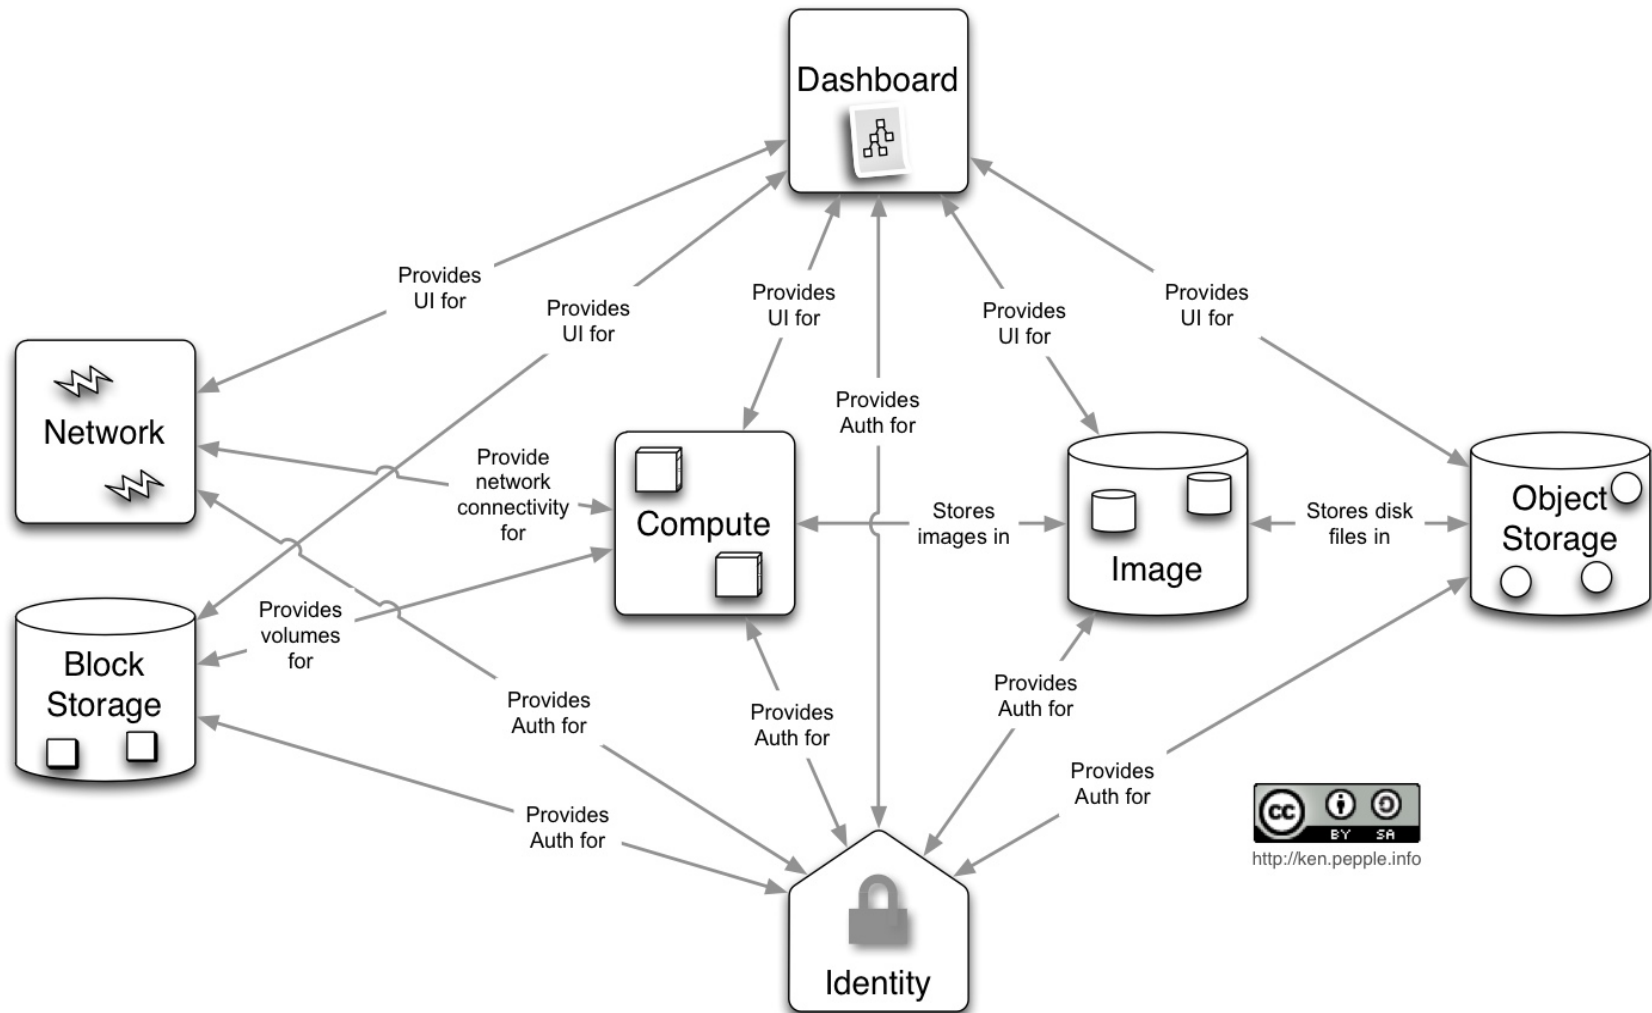


Network (Neutron) is a pluggable, scalable and API-driven system for managing networks and IP addresses. It allows users to customize network settings, then provision different network types (such as flat networks, VLANs or virtual private networks) on demand.




Block Storage (Cinder) provides persistent block storage guest to guest VMs. They're meant to be assigned to compute instances to allow for expanded storage. The Cinder software manages the creation of these blocks, plus the operations for attaching and detaching the blocks to compute servers.

OpenStack Core Services (Conceptual)



Neutron




Neutron Overview

- Incubation project in April 2011
- Promoted to Core Project at Folsom Summit (April 2012)
- Neutron solves two main problems in Nova – Network:
 1. Limited networking technology
 - Basic linux bridging-based implementation
 - Limited features (ACL, QoS, VPN?)
 - Limited multi-tenancy isolation – 802.1q VLAN tags
 2. Limited User/Tenant control over the network
 - Tenant cannot create their own network topologies
 - Tenant cannot leverage different network virtualization technologies


Why use Neutron?

- Provides REST APIs to manage network connections for the resources managed by other OpenStack services (e.g. Nova)
- Create rich topologies
- Technology agnostic
- Extensible
- Multi-tenancy: Isolation, Abstraction, full control over virtual networks
- Modular design: API specifies service, vendor provides its implementation
- Advance Services Support (LB, VPN, Firewall)

What does the user see?



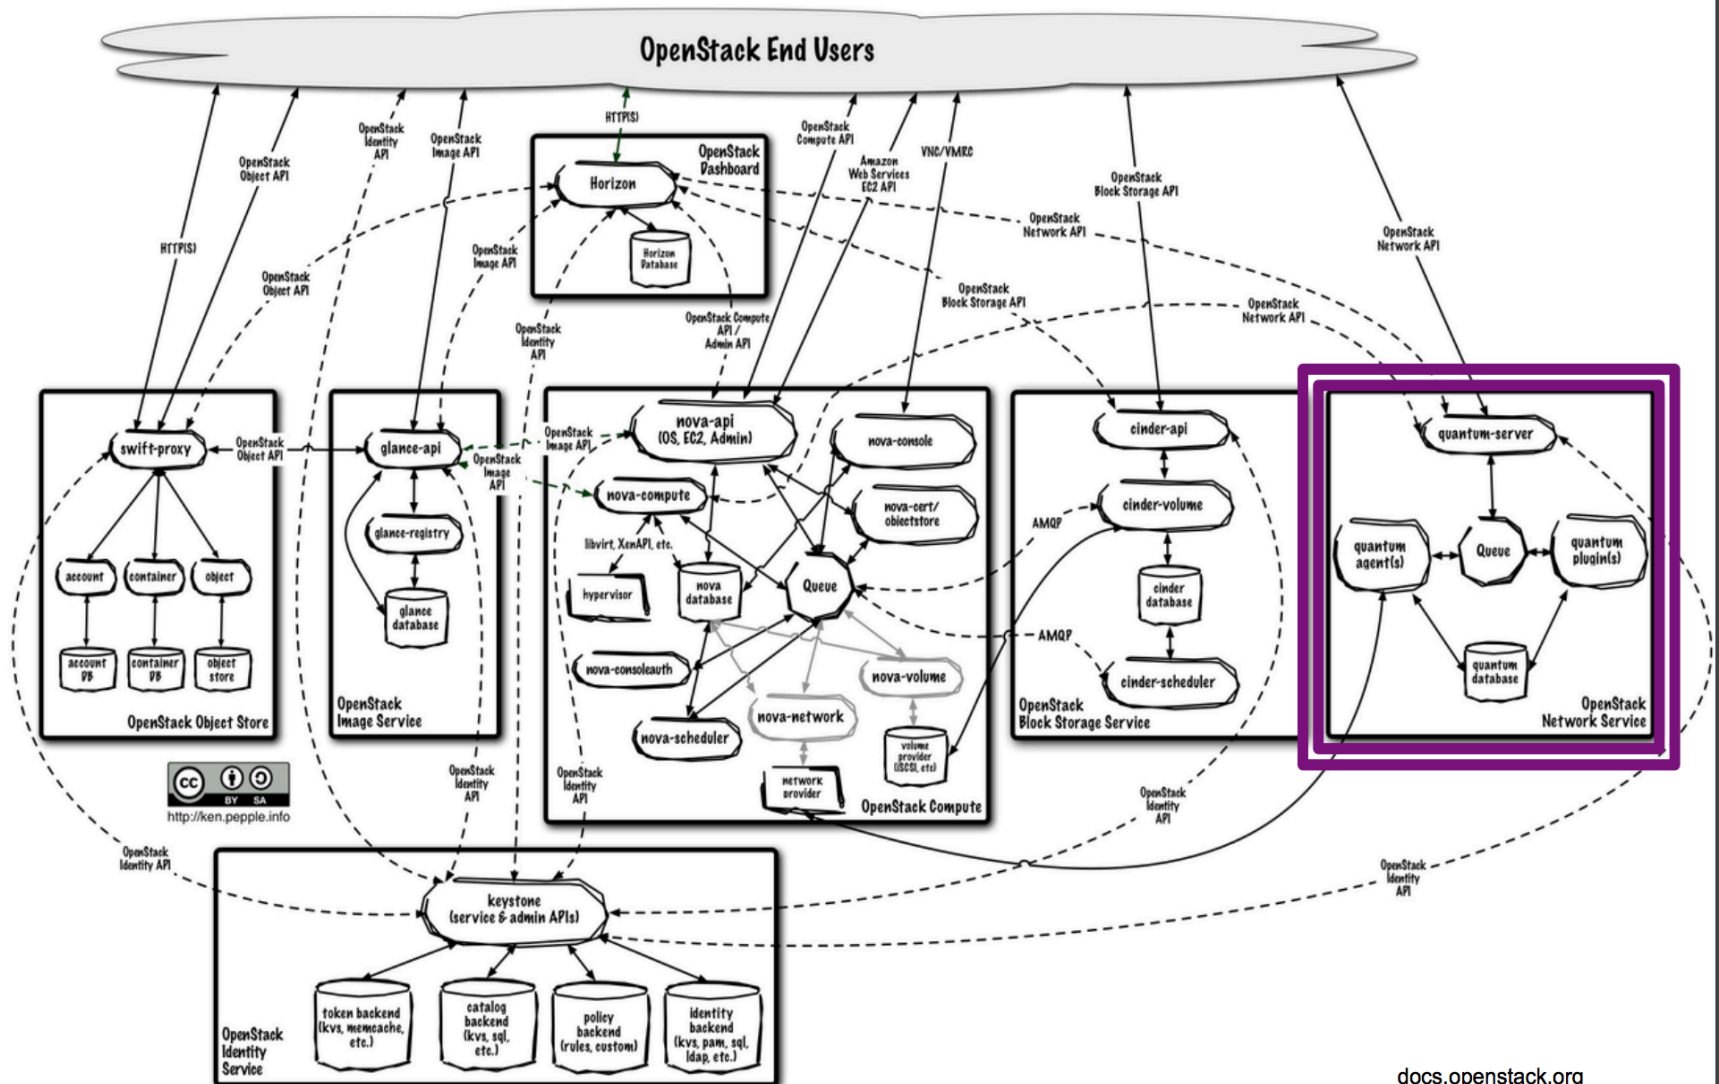
Abstractions



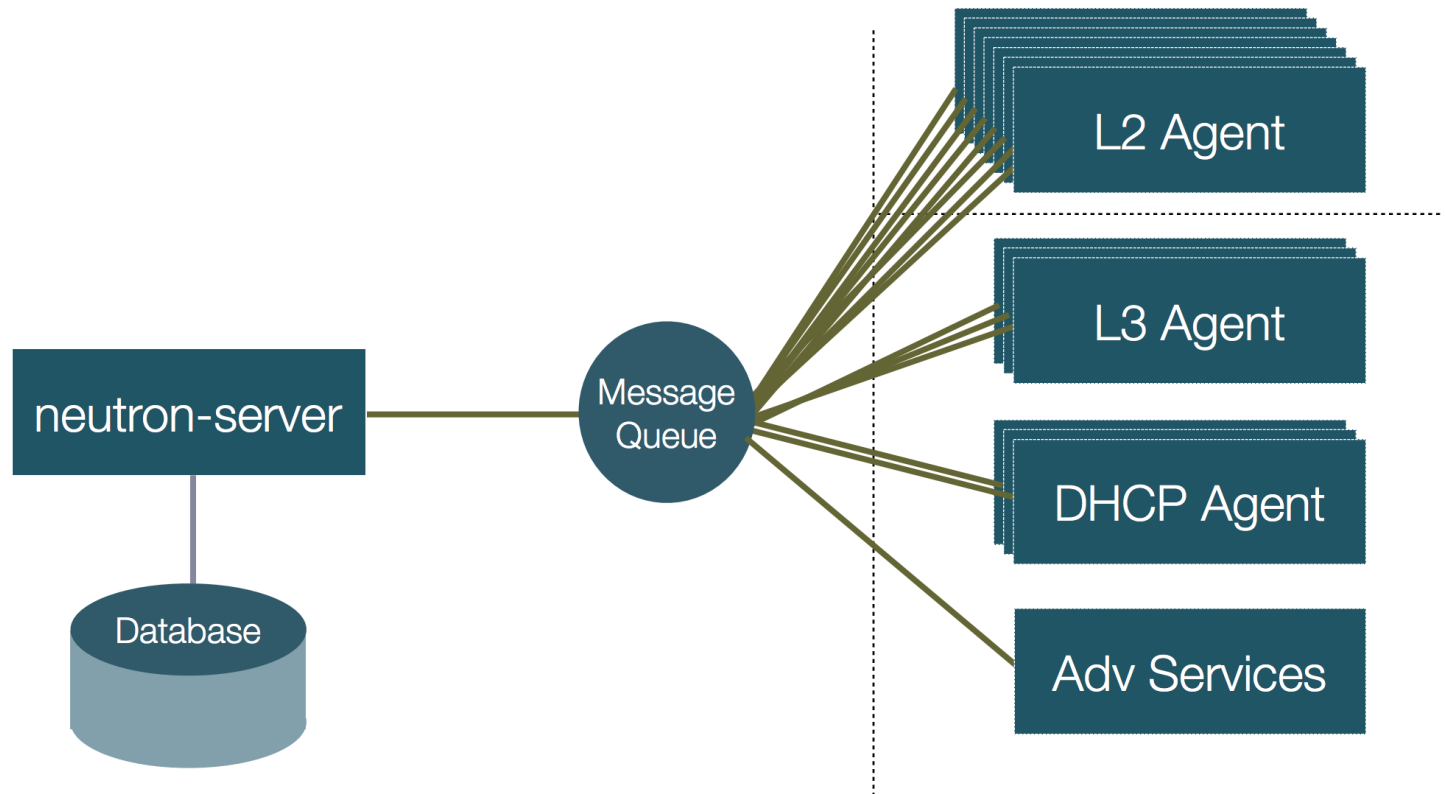
Design Goals

- Unified API
- Small Core
- Pluggable Open Architecture
- Extensible

Architecture




Basic Deployment




neutron-server

- REST API
 - HTTP(S) Python WSGI Application
 - Customary TCP port is 9696
 - Exposes logical resources
 - networks
 - subnets
 - ports
- Request/Response Serialization



neutron-server

- RPC Service
 - AMQP via Oslo messaging modules
 - Enables bidirectional agent communication
 - Optional



neutron-server

- PLUGIN
 - Written in Python
 - Only one is active
 - Must implement V2 API calls
 - Optional database access
 - Optional extension support



The Plugin

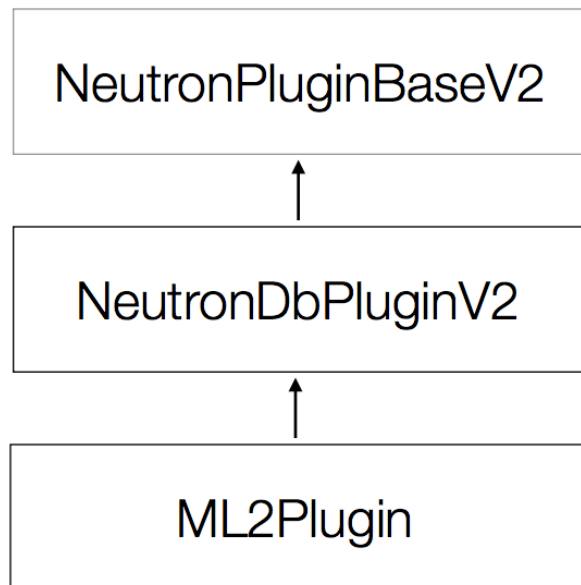
```
core_plugin =  
neutron.plugins.ml2.plugin.ML2Plugin
```

ML2Plugin

```
[DEFAULT]  
nova_admin_auth_url = http://192.168.57.10:35357/v2.0  
nova_admin_tenant_id = b9b4abce9d7d438783211363640dedda  
nova_admin_password = password  
nova_admin_username = nova  
nova_url = http://192.168.57.10:8774/v2  
notify_nova_on_port_data_changes = True  
notify_nova_on_port_status_change = True  
auth_strategy = keystone  
allow_overlapping_ips = True  
policy_file = /etc/neutron/policy.json  
debug = True  
verbose = True  
service_plugins = neutron.services.l3_router.l3_router_plugin.L3RouterPlugin  
core_plugin = neutron.plugins.ml2.plugin.ML2Plugin  
rabbit_password = yetanothersecret  
rabbit_hosts = 192.168.57.10  
rpc_backend = neutron.openstack.common.rpc.impl_kombu  
logging_exception_prefix = %(color)s%(asctime)s.%(msecs)03d TRACE %(name)s ^[[01  
;35m%(instance)s^[[00m  
logging_debug_format_suffix = ^[[00;33mfrom (pid=%(process)d) %(funcName)s %(pat  
hname)s:%(lineno)d^[[00m  
logging_default_format_string = %(asctime)s.%(msecs)03d %(color)s%(levelname)s %  
(name)s [^[[00;36m-%(color)s] ^[[01;35m%(instance)s%(color)s%(message)s^[[00m  
logging_context_format_string = %(asctime)s.%(msecs)03d %(color)s%(levelname)s %  
(name)s [^[[01;36m%(request_id)s ^[[00;36m%(user_name)s %(project_id)s%(color)s]  
^[[01;35m%(instance)s%(color)s%(message)s^[[00m  
state_path = /opt/stack/data/neutron  
# Print more verbose output (set logging level to INFO instead of default WARNIN  
G level).  
# verbose = False  
  
# Print debugging output (set logging level to DEBUG instead of default WARNING  
level).  
# debug = False
```

The Plugin

```
core_plugin =  
neutron.plugins.ml2.plugin.ML2Plugin
```



```
"""  
v2 Neutron Plug-in API specification.  
  
:class:`NeutronPluginBaseV2` provides the definition of minimum set of  
methods that needs to be implemented by a v2 Neutron Plug-in.  
"""  
  
import abc  
import six  
  
@six.add_metaclass(abc.ABCMeta)  
class NeutronPluginBaseV2(object):  
  
    @abc.abstractmethod  
    def create_subnet(self, context, subnet):  
        """Create a subnet.  
  
        Create a subnet, which represents a range of IP addresses  
        that can be allocated to devices  
  
        :param context: neutron api request context  
        :param subnet: dictionary describing the subnet, with keys  
                        as listed in the :obj:`RESOURCE_ATTRIBUTE_MAP` object  
                        in :file:`neutron/api/v2/attributes.py`. All keys will  
                        be populated.  
  
        """  
  
        pass  
  
    @abc.abstractmethod  
    def update_subnet(self, context, id, subnet):  
        """Update values of a subnet.  
  
        :param context: neutron api request context  
        :param id: UUID representing the subnet to update.  
        :param subnet: dictionary with keys indicating fields to update.  
                        valid keys are those that have a value of True for
```

Plugin Extensions

- Add logical resources to the REST API
- Discovered by server at startup:
 - REST: /v2.0/extensions
- Common extensions:
 - Binding, DHCP, L3, Provider, Quota, Security Groups
- Other extensions:
 - Extra Routes, Metering

Booting a VM

- nova boot

```
$ nova boot --flavor 1 --image cirros-0.3.1-x86_64-uec myinstance
```

+-----+-----+	
-----+	
Property	Value
+-----+-----+	
-----+	
OS-DCF:diskConfig	MANUAL
OS-EXT-AZ:availability_zone	nova
OS-EXT-STS:power_state	0

Booting a VM

- nova boot
- create port
 - notify DHCP of new port

```
+-----+
| id | fixed_ips |
+-----+
| 2eabf01c-5de9-48ad-8124-4cb26a4f106a | {"subnet_id": "4c4c957f-c816-4f68-9ca3-2efa283b5166", "ip_address": "192.168.123.2"} |
| 2f70f8b1-f10f-4152-a6b8-3395abc2fc87 | {"subnet_id": "4c4c957f-c816-4f68-9ca3-2efa283b5166", "ip_address": "192.168.123.4"} |
| 924dd91b-1cfc-4900-af6b-4977117640a7 | {"subnet_id": "4c4c957f-c816-4f68-9ca3-2efa283b5166", "ip_address": "192.168.123.3"} |
| c48dc021-cfd5-4f73-a6de-686404164f40 | {"subnet_id": "4c4c957f-c816-4f68-9ca3-2efa283b5166", "ip_address": "192.168.123.1"} |
+-----+
```

```
$ sudo ip netns list
qdhcp-01854654-cc1c-4549-8a9b-c579bbcc6e57
qrouter-4911f6bd-f389-44c4-824e-e8df867e5155
```

```
$ sudo ip netns exec qdhcp-01854654-cc1c-4549-8a9b-c579bbcc6e57 ip link list
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
13: tap924dd91b-1c: <BROADCAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UNKNOWN mode DEFAULT group default
   link/ether fa:16:3e:06:48:ea brd ff:ff:ff:ff:ff:ff
```

```
$ cat /opt/stack/data/neutron/dhcp/01854654-cc1c-4549-8a9b-c579bbcc6e57/opts
tag:tag0,option:router,192.168.123.1$ cat /opt/stack/data/neutron/dhcp/01854654-cc1c-4549-8a9b-c579bbcc6e57/host
fa:16:3e:b3:f7:08,host-192-168-123-2.openstacklocal,192.168.123.2
fa:16:3e:e3:c6:c6,host-192-168-123-1.openstacklocal,192.168.123.1
fa:16:3e:54:05:f9,host-192-168-123-4.openstacklocal,192.168.123.4
```

Booting a VM

- nova boot
- create port
 - notify DHCP of new port
- create device
- wire port
- boot

```
17: tap2f70f8b1-f1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast master qbr2f70f8b1-f1 state UNKNOWN  
mode DEFAULT group default qlen 500  
    link/ether fe:16:3e:54:05:f9 brd ff:ff:ff:ff:ff:ff
```

Summary

- Unified API
- Small Core
- Pluggable Open Architecture
 - Multiple Vendor Support
- Extensible

vmware®



Ryu OpenFlow
Controller

NEC



tail-f

em)brane)

ARISTA



I've Got It!



Acknowledgements

- www.docs.openstack.org
- <http://developer.openstack.org/api-ref-networking-v2.html>
- <http://docs.openstack.org/admin-guide-cloud/index.html>
- <http://www.slideshare.net/markmcclain/inside-architecture-of-neutron>